

# Fourth Down

Sam Hughes

1/8/2022

With this code, I created models for three occurrences vital in a fourth down calculator when a team chooses to go for it: conversion percentage, yards gained on a successful conversion attempt, and yards gained (or lost) on a failed conversion attempt. The purpose of this file is to show my findings along with cool data visualizations, so I will mute the cells unless they express important information about a model or a plot.

Objective 1: Use the data to make a logistic model to predict the conversion percentage when given a down and distance. Modeling conversion percentage will be difficult because the samples are relatively small when considering the amount of combinations between yards for first down and yards away from a touchdown. Therefore, I will make a few different models and choose the one that seems to perform the best via cross validation.

```
# First simpler model
gfi_model = glm(conversion ~ ydstogo + yardline, data = gfi_train, family = "binomial")
summary(gfi_model)
```

```
##
## Call:
## glm(formula = conversion ~ ydstogo + yardline, family = "binomial",
##      data = gfi_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.5528  -1.1903  -0.2801   1.0609   2.5560
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.363148   0.061542   5.901 3.62e-09 ***
## ydstogo      -0.145569   0.008220 -17.709 < 2e-16 ***
## yardline      0.007182   0.001463   4.908 9.20e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6401.8  on 4617  degrees of freedom
## Residual deviance: 6007.2  on 4615  degrees of freedom
## AIC: 6013.2
##
## Number of Fisher Scoring iterations: 3
```

```

gfi_model_intercept = gfi_model$coefficients[1]
gfi_model_coef_a = gfi_model$coefficients[2]
gfi_model_coef_b = gfi_model$coefficients[3]

ydstogo_and_yardline = function(ydstogo, yardline) {
  exp(gfi_model_intercept + gfi_model_coef_a * ydstogo + gfi_model_coef_b * yardline)/(1 + exp(gfi_model_intercept + gfi_model_coef_a * ydstogo + gfi_model_coef_b * yardline))
}
ydstogo_and_yardline(1, 15)

```

```

## (Intercept)
## 0.5806184

```

```

# Score model
gfi_test$prediction1 = ydstogo_and_yardline(gfi_test$ydstogo, gfi_test$yardline)
gfi_test$binary1 = 0
gfi_test[which(gfi_test$prediction1 >= 0.5), "binary1"] = 1
gfi_test$binary1 = as.factor(gfi_test$binary1)
gfi_test$conversion = as.factor(gfi_test$conversion)
str(gfi_test)

```

```

## tibble [1,154 x 6] (S3: tbl_df/tbl/data.frame)
## $ yardline      : num [1:1154] 69 52 86 43 38 45 31 39 6 31 ...
## $ ydstogo       : num [1:1154] 10 15 9 10 1 2 11 10 1 11 ...
## $ yards_gained  : num [1:1154] 9 0 0 30 2 3 31 8 0 -8 ...
## $ conversion    : Factor w/ 2 levels "0","1": 1 1 1 2 2 2 2 1 1 1 ...
## $ prediction1   : num [1:1154] 0.355 0.19 0.418 0.314 0.62 ...
## $ binary1       : Factor w/ 2 levels "0","1": 1 1 1 1 2 2 1 1 2 1 ...

```

```

sum(gfi_test$conversion == gfi_test$binary1)/nrow(gfi_test)

```

```

## [1] 0.6247834

```

```

confusionMatrix(gfi_test$binary1, gfi_test$conversion)

```

```

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##              0 271 145
##              1 288 450
##
##              Accuracy : 0.6248
##              95% CI : (0.5961, 0.6528)
##              No Information Rate : 0.5156
##              P-Value [Acc > NIR] : 5.094e-14
##
##              Kappa : 0.243
##
##              Mcnemar's Test P-Value : 8.849e-12
##
##              Sensitivity : 0.4848

```

```
##           Specificity : 0.7563
##           Pos Pred Value : 0.6514
##           Neg Pred Value : 0.6098
##           Prevalence : 0.4844
##           Detection Rate : 0.2348
##           Detection Prevalence : 0.3605
##           Balanced Accuracy : 0.6205
##
##           'Positive' Class : 0
##
```

```
# second model with squared and cubed terms for ydstogo
```

```
gfi_model_2 = glm(conversion ~ ydstogo + ydstogo_squared + ydstogo_cubed + yardline, data = gfi_train,
summary(gfi_model_2)
```

```
##
## Call:
## glm(formula = conversion ~ ydstogo + ydstogo_squared + ydstogo_cubed +
##       yardline, family = "binomial", data = gfi_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6222  -1.1170  -0.2285   1.0365   2.5298
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.6566396   0.0854179   7.687 1.50e-14 ***
## ydstogo        -0.3445208   0.0440933  -7.813 5.56e-15 ***
## ydstogo_squared  0.0216148   0.0055968   3.862 0.000112 ***
## ydstogo_cubed   -0.0005438   0.0001847  -2.945 0.003234 **
## yardline        0.0076158   0.0014713   5.176 2.27e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6401.8  on 4617  degrees of freedom
## Residual deviance: 5978.3  on 4613  degrees of freedom
## AIC: 5988.3
##
## Number of Fisher Scoring iterations: 6
```

```
gfi_model_intercept2 = gfi_model_2$coefficients[1]
gfi_model_coef_a2 = gfi_model_2$coefficients[2]
gfi_model_coef_b2 = gfi_model_2$coefficients[3]
gfi_model_coef_c2 = gfi_model_2$coefficients[4]
gfi_model_coef_d2 = gfi_model_2$coefficients[5]

ydstogo_and_yardline_2 = function(ydstogo, yardline) {
  exp(gfi_model_intercept2 + gfi_model_coef_a2 * ydstogo + gfi_model_coef_b2 * ydstogo^2 + gfi_model_coef_c2 * ydstogo^3 + gfi_model_coef_d2 * yardline)
}

ydstogo_and_yardline_2(1, 15)
```

```
## (Intercept)
## 0.6100271
```

```
# Score model
```

```
gfi_test$prediction2 = ydstogo_and_yardline_2(gfi_test$ydstogo, gfi_test$yardline)
gfi_test$binary2 = 0
gfi_test[which(gfi_test$prediction2 >= 0.5), "binary2"] = 1
gfi_test$binary2 = as.factor(gfi_test$binary2)
gfi_test$conversion = as.factor(gfi_test$conversion)
str(gfi_test)
```

```
## tibble [1,154 x 8] (S3: tbl_df/tbl/data.frame)
## $ yardline      : num [1:1154] 69 52 86 43 38 45 31 39 6 31 ...
## $ ydstogo       : num [1:1154] 10 15 9 10 1 2 11 10 1 11 ...
## $ yards_gained  : num [1:1154] 9 0 0 30 2 3 31 8 0 -8 ...
## $ conversion    : Factor w/ 2 levels "0","1": 1 1 1 2 2 2 2 1 1 1 ...
## $ prediction1   : num [1:1154] 0.355 0.19 0.418 0.314 0.62 ...
## $ binary1       : Factor w/ 2 levels "0","1": 1 1 1 1 2 2 1 1 2 1 ...
## $ prediction2   : num [1:1154] 0.344 0.252 0.393 0.301 0.651 ...
## $ binary2       : Factor w/ 2 levels "0","1": 1 1 1 1 2 2 1 1 2 1 ...
```

```
sum(gfi_test$conversion == gfi_test$binary2)/nrow(gfi_test)
```

```
## [1] 0.6334489
```

```
confusionMatrix(gfi_test$binary2, gfi_test$conversion)
```

```
## Confusion Matrix and Statistics
```

```
##
##              Reference
## Prediction    0    1
##              0 304 168
##              1 255 427
##
##              Accuracy : 0.6334
##              95% CI   : (0.6049, 0.6613)
##              No Information Rate : 0.5156
##              P-Value [Acc > NIR] : 4.560e-16
##
##              Kappa   : 0.2627
##
##              McNemar's Test P-Value : 2.896e-05
##
##              Sensitivity : 0.5438
##              Specificity : 0.7176
##              Pos Pred Value : 0.6441
##              Neg Pred Value : 0.6261
##              Prevalence   : 0.4844
##              Detection Rate : 0.2634
##              Detection Prevalence : 0.4090
##              Balanced Accuracy : 0.6307
##
```

```
##      'Positive' Class : 0
##
```

```
gfi_model_3 = glm(conversion ~ ydstogo_ln + yardline_ln, data = gfi_train, family = "binomial")
summary(gfi_model_3)
```

```
##
## Call:
## glm(formula = conversion ~ ydstogo_ln + yardline_ln, family = "binomial",
##      data = gfi_train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.5809  -1.0802  -0.6485   1.0629   1.8549
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.15348    0.09272   1.655  0.0979 .
## ydstogo_ln  -0.68594    0.03474 -19.747 < 2e-16 ***
## yardline_ln  0.16939    0.02873   5.896 3.72e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 6401.8  on 4617  degrees of freedom
## Residual deviance: 5975.3  on 4615  degrees of freedom
## AIC: 5981.3
##
## Number of Fisher Scoring iterations: 4
```

```
ydstogo_and_yardline_3 = function(ydstogo, yardline) {
  exp(gfi_model_3$coefficients[1] + gfi_model_3$coefficients[2] * log(ydstogo) + gfi_model_3$coefficients[3] * log(yardline))
}
ydstogo_and_yardline_3(1, 15)
```

```
## (Intercept)
##      0.6484413
```

```
# Score model
gfi_test$prediction3 = ydstogo_and_yardline_3(gfi_test$ydstogo, gfi_test$yardline)
gfi_test$binary3 = 0
gfi_test[which(gfi_test$prediction3 >= 0.5), "binary3"] = 1
gfi_test$binary3 = as.factor(gfi_test$binary3)
gfi_test$conversion = as.factor(gfi_test$conversion)
str(gfi_test)
```

```
## tibble [1,154 x 10] (S3: tbl_df/tbl/data.frame)
## $ yardline      : num [1:1154] 69 52 86 43 38 45 31 39 6 31 ...
## $ ydstogo       : num [1:1154] 10 15 9 10 1 2 11 10 1 11 ...
## $ yards_gained : num [1:1154] 9 0 0 30 2 3 31 8 0 -8 ...
## $ conversion    : Factor w/ 2 levels "0","1": 1 1 1 2 2 2 2 1 1 1 ...
```

```
## $ prediction1 : num [1:1154] 0.355 0.19 0.418 0.314 0.62 ...
## $ binary1      : Factor w/ 2 levels "0","1": 1 1 1 1 2 2 1 1 2 1 ...
## $ prediction2 : num [1:1154] 0.344 0.252 0.393 0.301 0.651 ...
## $ binary2      : Factor w/ 2 levels "0","1": 1 1 1 1 2 2 1 1 2 1 ...
## $ prediction3 : num [1:1154] 0.33 0.262 0.355 0.312 0.683 ...
## $ binary3      : Factor w/ 2 levels "0","1": 1 1 1 1 2 2 1 1 2 1 ...
```

```
sum(gfi_test$conversion == gfi_test$binary3)/nrow(gfi_test)
```

```
## [1] 0.6360485
```

```
confusionMatrix(gfi_test$binary3, gfi_test$conversion)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  0    1
##           0 332 193
##           1 227 402
##
##           Accuracy : 0.636
##           95% CI : (0.6075, 0.6639)
##           No Information Rate : 0.5156
##           P-Value [Acc > NIR] : <2e-16
##
##           Kappa : 0.27
##
## Mcnemar's Test P-Value : 0.1073
##
##           Sensitivity : 0.5939
##           Specificity : 0.6756
##           Pos Pred Value : 0.6324
##           Neg Pred Value : 0.6391
##           Prevalence : 0.4844
##           Detection Rate : 0.2877
##           Detection Prevalence : 0.4549
##           Balanced Accuracy : 0.6348
##
##           'Positive' Class : 0
##
```


The confusion matrices give us information against the testing data, and we can also visualize this data via ggplot.

```
ggplot(data = subset(gfi_test_agg, ydstogo == 1),
       mapping = aes(x = yardline,
                     y = mean))+
  ggtitle("4th and 1 Conversion by Yards to Endzone", subtitle = "2010 to 2020, Test Data")+
  xlab("Yards to Endzone")+
  ylab("Conversion Percentage")+
  geom_point(aes(size = count))+
  geom_line(aes(y = ydstogo_and_yardline(ydstogo, yardline),
```

```

    size = 7, colour = "Simple Model"))+
  geom_line(aes(y = ydstogo_and_yardline_2(ydstogo, yardline),
    size = 7, colour = "Complex Model"))+
  geom_line(aes(y = ydstogo_and_yardline_3(ydstogo, yardline),
    size = 7, colour = "Ln Model"))

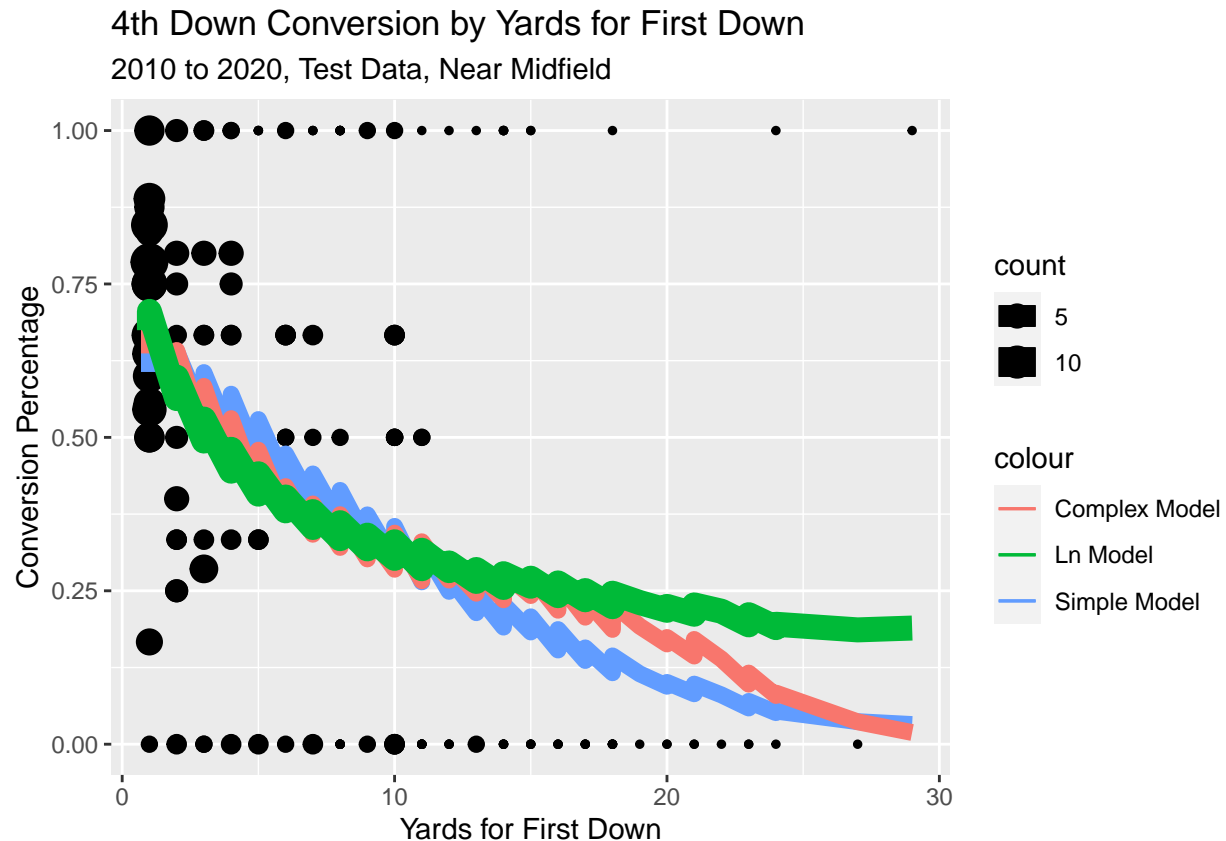
```



```

ggplot(data = subset(gfi_test_agg, yardline >= 30 & yardline <= 70),
  mapping = aes(x = ydstogo,
    y = mean))+
  ggtitle("4th Down Conversion by Yards for First Down",
    subtitle = "2010 to 2020, Test Data, Near Midfield")+
  xlab("Yards for First Down")+
  ylab("Conversion Percentage")+
  geom_point(aes(size = count))+
  geom_line(aes(y = ydstogo_and_yardline(ydstogo, yardline),
    size = 3, colour = "Simple Model"))+
  geom_line(aes(y = ydstogo_and_yardline_2(ydstogo, yardline),
    size = 3, colour = "Complex Model"))+
  geom_line(aes(y = ydstogo_and_yardline_3(ydstogo, yardline),
    size = 7, colour = "Ln Model"))

```



I think the logarithmic model is the best choice because of its performance at short yardage (when going for it is a realistic possibility). Fourth and 1 mean percentage  $\rightarrow$  0.65, 2  $\rightarrow$  0.56, 3  $\rightarrow$  0.49, all of which the logarithmic model is closest to predicting

Objective 2: When a conversion does happen, find out the excess yards gained on the play (the yards gained beyond the first down marker). Let's start with a simple model of just the distance for a first down and distance for a touchdown.

```
# while converting
gfi_success = subset(df_fourth_go_simple, conversion == 1)
gfi_success$excess = gfi_success$yards_gained - gfi_success$ydstogo

# Split data
set.seed(123)
training_samples_gfi = createDataPartition(gfi_success$excess, p = 0.8, list = FALSE)
gfi_train = gfi_success[training_samples_gfi, ]
gfi_test = gfi_success[-training_samples_gfi, ]

# for graph later
gfi_test_agg = gfi_test %>%
  group_by(ydstogo, yardline) %>%
  summarise("mean" = mean(excess), "count" = length(excess))
```

## 'summarise()' has grouped output by 'ydstogo'. You can override using the '.groups' argument.



```

# for more complex models later
gfi_train$ydstogo_squared = gfi_train$ydstogo^2
gfi_train$ydstogo_cubed = gfi_train$ydstogo^3

gfi_train$yardline_squared = gfi_train$yardline^2
gfi_train$yardline_cubed = gfi_train$yardline^3

# simple model
summary(lm(excess ~ ydstogo * yardline, data = gfi_train))

##
## Call:
## lm(formula = excess ~ ydstogo * yardline, data = gfi_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -17.251  -4.871  -2.329   1.654   56.503
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    1.654237    0.470890   3.513 0.000452 ***
## ydstogo         0.424376    0.118992   3.566 0.000369 ***
## yardline        0.098859    0.011158   8.860 < 2e-16 ***
## ydstogo:yardline -0.001608    0.002242  -0.717 0.473190
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.58 on 2313 degrees of freedom
## Multiple R-squared:  0.08653,    Adjusted R-squared:  0.08535
## F-statistic: 73.04 on 3 and 2313 DF,  p-value: < 2.2e-16

anova(lm(excess ~ ydstogo * yardline, data = gfi_train))

## Analysis of Variance Table
##
## Response: excess
##              Df Sum Sq Mean Sq  F value Pr(>F)
## ydstogo         1    7238   7238.3   98.3166 <2e-16 ***
## yardline         1   8855   8855.0  120.2758 <2e-16 ***
## ydstogo:yardline  1     38    37.9    0.5147 0.4732
## Residuals      2313 170289    73.6
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

# interaction term is insignificant
summary(lm(excess ~ ydstogo + yardline, data = gfi_train))

##
## Call:
## lm(formula = excess ~ ydstogo + yardline, data = gfi_train)
##

```

```
## Residuals:
##      Min       1Q   Median       3Q      Max
## -17.849  -4.873  -2.342   1.722  56.597
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.870536    0.361676   5.172 2.52e-07 ***
## ydstogo      0.347415    0.051483   6.748 1.89e-11 ***
## yardline     0.093711    0.008544  10.968 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.579 on 2314 degrees of freedom
## Multiple R-squared:  0.08633,    Adjusted R-squared:  0.08554
## F-statistic: 109.3 on 2 and 2314 DF,  p-value: < 2.2e-16
```

```
anova(lm(excess ~ ydstogo + yardline, data = gfi_train))
```

```
## Analysis of Variance Table
##
## Response: excess
##              Df Sum Sq Mean Sq F value    Pr(>F)
## ydstogo        1   7238   7238.3   98.337 < 2.2e-16 ***
## yardline        1   8855   8855.0  120.301 < 2.2e-16 ***
## Residuals    2314 170327    73.6
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# looks solid
```

```
excess_model1 = lm(excess ~ ydstogo + yardline, data = gfi_train)
excess_model1_intercept = excess_model1$coefficients[1]
excess_model1_coef_a = excess_model1$coefficients[2]
excess_model1_coef_b = excess_model1$coefficients[3]

excess_model1_equation = function(ydstogo, yardline) {
  excess_model1_intercept + excess_model1_coef_a * ydstogo + excess_model1_coef_b * yardline
}

excess_model1_equation(5, 10)
```

```
## (Intercept)
##      4.544718
```

```
excess_model1_equation(30, 90)
```

```
## (Intercept)
##      20.72696
```

```
# check performance with test data
```

```
gfi_test$prediction1 = excess_model1_equation(gfi_test$ydstogo, gfi_test$yardline)
gfi_test$error1 = abs(gfi_test$excess - gfi_test$prediction1)
```

```
# MAE
sum(gfi_test$error1)/nrow(gfi_test)
```

```
## [1] 5.74256
```

```
# more advanced model
```

```
summary(lm(excess ~ ydstogo + ydstogo_squared + ydstogo_cubed + yardline + yardline_squared + yardline_cubed, data = gfi_train))
```

```
##
## Call:
## lm(formula = excess ~ ydstogo + ydstogo_squared + ydstogo_cubed +
##     yardline + yardline_squared + yardline_cubed, data = gfi_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -17.247  -5.172  -2.361   1.426  58.413
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -2.621e+00  6.149e-01  -4.263 2.10e-05 ***
## ydstogo        1.721e+00  2.505e-01   6.873 8.08e-12 ***
## ydstogo_squared -1.528e-01  3.124e-02  -4.891 1.07e-06 ***
## ydstogo_cubed   3.724e-03  9.660e-04   3.856 0.000119 ***
## yardline       3.836e-01  6.029e-02   6.362 2.39e-10 ***
## yardline_squared -6.627e-03  1.770e-03  -3.745 0.000185 ***
## yardline_cubed  3.891e-05  1.434e-05   2.713 0.006725 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 8.423 on 2310 degrees of freedom
## Multiple R-squared:  0.1209, Adjusted R-squared:  0.1187
## F-statistic: 52.97 on 6 and 2310 DF, p-value: < 2.2e-16
```

```
anova(lm(excess ~ ydstogo + ydstogo_squared + ydstogo_cubed + yardline + yardline_squared + yardline_cubed, data = gfi_train))
```

```
## Analysis of Variance Table
##
## Response: excess
##              Df Sum Sq Mean Sq  F value    Pr(>F)
## ydstogo         1    7238   7238.3 102.0328 < 2.2e-16 ***
## ydstogo_squared  1    2404   2403.6  33.8814 6.672e-09 ***
## ydstogo_cubed   1    1260   1260.1  17.7619 2.600e-05 ***
## yardline        1    8678   8678.1 122.3280 < 2.2e-16 ***
## yardline_squared  1    2444   2444.1  34.4531 4.995e-09 ***
## yardline_cubed   1     522    522.0   7.3582 0.006725 **
## Residuals     2310 163874    70.9
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
# every single term here is significant

excess_model2 = lm(excess ~ ydstogo + ydstogo_squared + ydstogo_cubed + yardline + yardline_squared + y

excess_model2_equation = function(ydstogo, yardline) {
  excess_model2$coefficients[1] + excess_model2$coefficients[2] * ydstogo + excess_model2$coefficients[
}

excess_model2_equation(1, 1)
```

```
## (Intercept)
## -0.6718014
```

```
excess_model2_equation(30, 90)
```

```
## (Intercept)
## 21.29527
```

```
gfi_test$prediction2 = excess_model2_equation(gfi_test$ydstogo, gfi_test$yardline)
gfi_test$error2 = abs(gfi_test$excess - gfi_test$prediction2)
```

```
# MAE
sum(gfi_test$error2)/nrow(gfi_test)
```

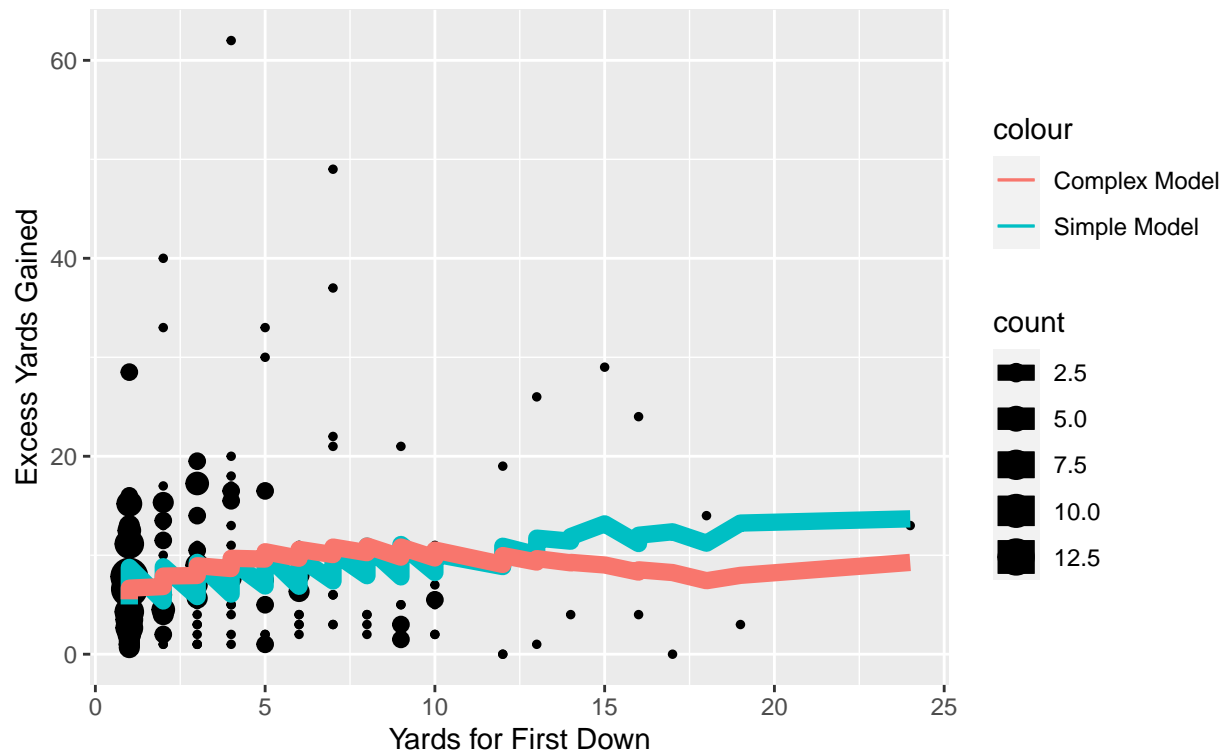
```
## [1] 5.641885
```

We can now visually compare the two models.

```
ggplot(data = subset(gfi_test_agg, yardline >= 30 & yardline <= 70),
  mapping = aes(x = ydstogo,
    y = mean))+
  ggtitle("Excess Yards Gained on 4th Down Conversions",
    subtitle = "2010 to 2020, Test Data, Near Midfield")+
  xlab("Yards for First Down")+
  ylab("Excess Yards Gained")+
  geom_point(aes(size = count))+
  geom_line(aes(y = excess_model1_equation(ydstogo, yardline),
    size = 3, colour = "Simple Model"))+
  geom_line(aes(y = excess_model2_equation(ydstogo, yardline),
    size = 3, colour = "Complex Model"))
```

## Excess Yards Gained on 4th Down Conversions

### 2010 to 2020, Test Data, Near Midfield



There does not seem to be much difference at all between the two equations. The MAE values are similar and the models are similar visually. I would say the simpler one is slightly better because the complex model sometimes gives negatives in niche situations which should be impossible.

Objective 3: The opposite of Objective 2. When teams fail to convert on fourth down, predict the amount of yards short they are from the first down marker.

```
# make simple model
summary(lm(shortage ~ ydstogo + yardline, data = gfi_train))

##
## Call:
## lm(formula = shortage ~ ydstogo + yardline, data = gfi_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -19.8721  -0.8952  -0.6129   0.0290  28.6493
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.950697   0.161827   5.875 4.85e-09 ***
## ydstogo      0.899811   0.015938  56.458 < 2e-16 ***
## yardline     0.002791   0.003850   0.725  0.469
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## Residual standard error: 3.948 on 2301 degrees of freedom
## Multiple R-squared:  0.6128, Adjusted R-squared:  0.6125
## F-statistic: 1821 on 2 and 2301 DF,  p-value: < 2.2e-16
```

```
anova(lm(shortage ~ ydstogo + yardline, data = gfi_train))
```

```
## Analysis of Variance Table
##
## Response: shortage
##           Df Sum Sq Mean Sq  F value Pr(>F)
## ydstogo      1  56764   56764 3640.9645 <2e-16 ***
## yardline      1     8         8   0.5256 0.4685
## Residuals 2301  35873     16
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

*# Yardline seems insignificant. Let's remove it*

```
summary(lm(shortage ~ ydstogo, data = gfi_train))
```

```
##
## Call:
## lm(formula = shortage ~ ydstogo, data = gfi_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -19.9118  -0.9329  -0.6441   0.0671  28.6070
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.02911     0.12036    8.55  <2e-16 ***
## ydstogo      0.90376     0.01498   60.35  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.948 on 2302 degrees of freedom
## Multiple R-squared:  0.6127, Adjusted R-squared:  0.6125
## F-statistic: 3642 on 1 and 2302 DF,  p-value: < 2.2e-16
```

```
anova(lm(shortage ~ ydstogo, data = gfi_train))
```

```
## Analysis of Variance Table
##
## Response: shortage
##           Df Sum Sq Mean Sq F value    Pr(>F)
## ydstogo      1  56764   56764 3641.7 < 2.2e-16 ***
## Residuals 2302  35881     16
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```

# Looks better now

shortage_model1 = lm(shortage ~ ydstogo, data = gfi_train)
shortage_model1_intercept = excess_model1$coefficients[1]
shortage_model1_coef_a = excess_model1$coefficients[2]

shortage_model1_equation = function(ydstogo) {
  shortage_model1_intercept + shortage_model1_coef_a * ydstogo
}

shortage_model1_equation(1)

## (Intercept)
##      2.217951

shortage_model1_equation(10)

## (Intercept)
##      5.344684

# check performance with test data

gfi_test$prediction1 = shortage_model1_equation(gfi_test$ydstogo)
gfi_test$error1 = abs(gfi_test$shortage - gfi_test$prediction1)

# MAE
sum(gfi_test$error1)/nrow(gfi_test)

## [1] 3.270582

# complex model
gfi_train$sq = gfi_train$ydstogo^2
gfi_train$cu = gfi_train$ydstogo^3

summary(lm(shortage ~ ydstogo + sq + cu, data = gfi_train))

##
## Call:
## lm(formula = shortage ~ ydstogo + sq + cu, data = gfi_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -19.1186  -0.7773  -0.6577  -0.0835   28.8444
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.5500514  0.1845917   2.980  0.00291 **
## ydstogo      1.1244930  0.0701341  16.033 < 2e-16 ***
## sq          -0.0171321  0.0061302  -2.795  0.00524 **
## cu           0.0002932  0.0001396   2.100  0.03586 *
## ---

```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.939 on 2300 degrees of freedom
## Multiple R-squared:  0.6147, Adjusted R-squared:  0.6142
## F-statistic: 1223 on 3 and 2300 DF,  p-value: < 2.2e-16
```

```
anova(lm(shortage ~ ydstogo + sq + cu, data = gfi_train))
```

```
## Analysis of Variance Table
```

```
##
```

```
## Response: shortage
```

```
##          Df Sum Sq Mean Sq  F value    Pr(>F)
## ydstogo    1  56764   56764 3657.7228 < 2.2e-16 ***
## sq         1   120     120   7.7100  0.005536 **
## cu         1    68     68   4.4088  0.035862 *
## Residuals 2300  35693     16
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
shortage_model2 = lm(shortage ~ ydstogo + sq + cu, data = gfi_train)
```

```
shortage_model2_intercept = shortage_model2$coefficients[1]
```

```
shortage_model2_coef_a = shortage_model2$coefficients[2]
```

```
shortage_model2_coef_b = shortage_model2$coefficients[3]
```

```
shortage_model2_coef_c = shortage_model2$coefficients[4]
```

```
shortage_model2_equation = function(ydstogo) {
```

```
  shortage_model2_intercept + shortage_model2_coef_a * ydstogo + shortage_model2_coef_b * ydstogo^2 + shortage_model2_coef_c * ydstogo^3
}
```

```
shortage_model2_equation(1)
```

```
## (Intercept)
```

```
##      1.657705
```

```
shortage_model2_equation(10)
```

```
## (Intercept)
```

```
##     10.37494
```

```
# check performance with test data
```

```
gfi_test$prediction2 = shortage_model2_equation(gfi_test$ydstogo)
```

```
gfi_test$error2 = abs(gfi_test$shortage - gfi_test$prediction2)
```

```
# MAE
```

```
sum(gfi_test$error2)/nrow(gfi_test)
```

```
## [1] 1.816694
```



```
# model with interaction term
summary(lm(shortage ~ ydstogo * yardline, data = gfi_train))
```

```
##
## Call:
## lm(formula = shortage ~ ydstogo * yardline, data = gfi_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -19.9528  -0.8842  -0.5949  -0.0306   28.5089
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    0.6168550   0.2172696    2.839  0.00456 **
## ydstogo         0.9680506   0.0336715   28.750 < 2e-16 ***
## yardline        0.0114675   0.0053873    2.129  0.03339 *
## ydstogo:yardline -0.0014922  0.0006487   -2.300  0.02153 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.945 on 2300 degrees of freedom
## Multiple R-squared:  0.6137, Adjusted R-squared:  0.6132
## F-statistic: 1218 on 3 and 2300 DF, p-value: < 2.2e-16
```

```
anova(lm(shortage ~ ydstogo * yardline, data = gfi_train))
```

```
## Analysis of Variance Table
##
## Response: shortage
##              Df Sum Sq Mean Sq  F value    Pr(>F)
## ydstogo         1  56764   56764 3647.7532 < 2e-16 ***
## yardline         1     8      8    0.5266 0.46811
## ydstogo:yardline 1    82     82    5.2903 0.02153 *
## Residuals      2300 35791     16
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
shortage_model3 = lm(shortage ~ ydstogo * yardline, data = gfi_train)
shortage_model3_intercept = shortage_model3$coefficients[1]
shortage_model3_coef_a = shortage_model3$coefficients[2]
shortage_model3_coef_b = shortage_model3$coefficients[3]
shortage_model3_coef_c = shortage_model3$coefficients[4]

shortage_model3_equation = function(ydstogo, yardline) {
  shortage_model3_intercept + shortage_model3_coef_a * ydstogo + shortage_model3_coef_b * yardline + sh
}

shortage_model3_equation(5, 10)
```

```
## (Intercept)
##      5.497175
```

```

gfi_test$prediction3 = shortage_model3_equation(gfi_test$ydstogo, gfi_test$yardline)
gfi_test$error3 = abs(gfi_test$shortage - gfi_test$prediction3)

# MAE
sum(gfi_test$error3)/nrow(gfi_test)

```

```
## [1] 1.797503
```

Visualition of Objective 3:

```

gfi_test_agg = gfi_test %>%
  group_by(ydstogo, yardline) %>%
  summarise("mean" = mean(shortage), "count" = length(shortage))

```

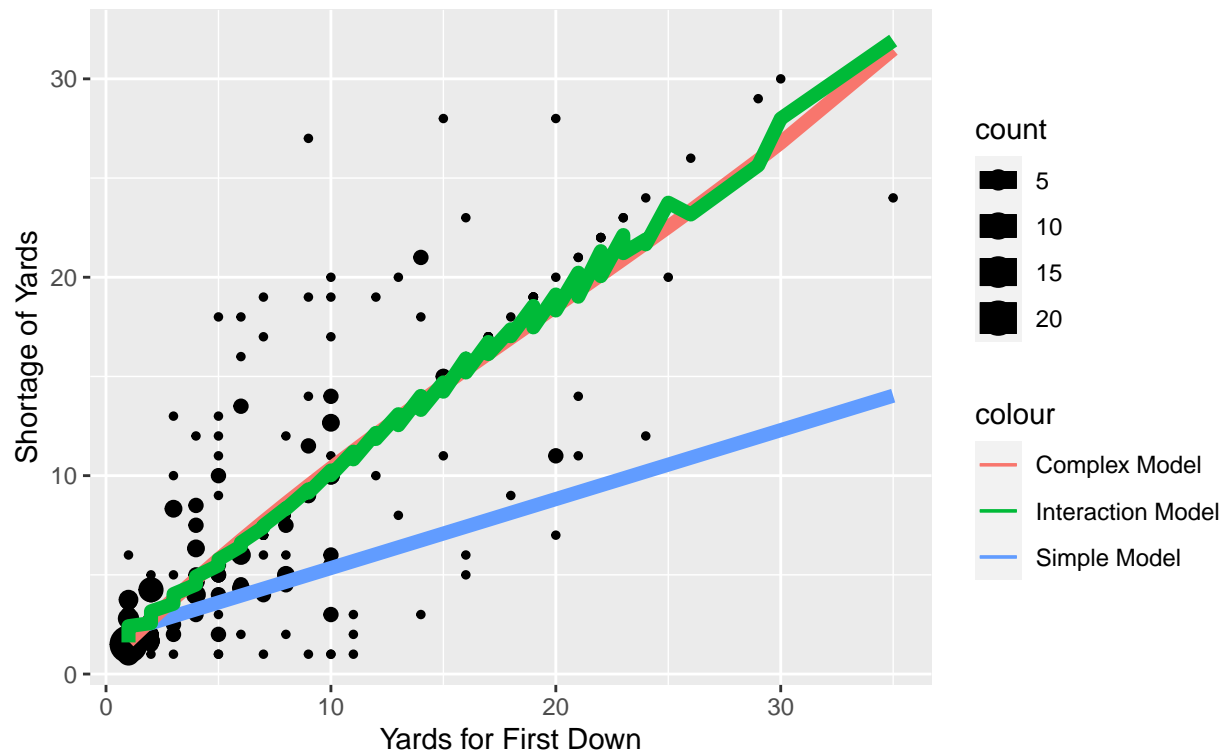
## 'summarise()' has grouped output by 'ydstogo'. You can override using the '.groups' argument.

```

ggplot(data = subset(gfi_test_agg, yardline >= 0 & yardline <= 100),
  mapping = aes(x = ydstogo,
    y = mean)) +
  ggtitle("Yards Short on Failed 4th Down Conversions",
    subtitle = "2010 to 2020, Test Data") +
  xlab("Yards for First Down") +
  ylab("Shortage of Yards") +
  geom_point(aes(size = count)) +
  geom_line(aes(y = shortage_model1_equation(ydstogo),
    size = 3, colour = "Simple Model")) +
  geom_line(aes(y = shortage_model2_equation(ydstogo),
    size = 3, colour = "Complex Model")) +
  geom_line(aes(y = shortage_model3_equation(ydstogo, yardline),
    size = 3, colour = "Interaction Model"))

```

Yards Short on Failed 4th Down Conversions  
2010 to 2020, Test Data



It seems as if either the interaction term or complex model are the best. I think the best bet is the interaction term model because it has the lowest MAE.